



Lab Assignment 3: Raft

Notes

- This lab is assigned to students at MIT.
- [You should work on this assignment individually.](#)
- This lab is due in three parts. You must submit each part on the corresponding due date. This lab does not involve a lot of code, but concurrency makes it potentially challenging to debug; start each part early.
- The deadline for part A is March 28 2017, for part B is April 18 2017, and for part C is April 25 2017.

Introduction

In this and next lab, you'll build a fault-tolerant key/value storage system. In this lab you'll implement Raft, a replicated state machine protocol. In the next lab you'll build a key/value service on top of Raft.

A replicated service (e.g., key/value database) achieves fault tolerance by storing copies of its data on multiple replica servers. Replication allows the service to continue operating even if some of its servers experience failures (crashes or a broken or flaky network). The challenge is that failures may cause the replicas to hold differing copies of the data.

Raft manages a service's state replicas, and in particular it helps the service sort out what the correct state is after failures. Raft implements a replicated state machine. It organizes client requests into a sequence, called the log, and ensures that all the replicas agree on the contents of the log. Each replica executes the client requests in the log in the order they appear in the log, applying those requests to the replica's local copy of the service's state. Since all the live replicas see the same log contents, they all execute the same requests in the same order, and thus continue to have identical service state. If a server fails but later recovers, Raft takes care of bringing its log up to date. Raft will continue to operate as long as at least a majority of the servers are alive and can talk to each other. If there is no such majority, Raft will make no progress, but will pick up where it left off as soon as a majority can communicate again.

In this lab you'll implement Raft as a Go object type with associated methods, meant to be used as a module in a larger service. A set of Raft instances talk to each other with RPC to maintain replicated logs. Your Raft interface will support an indefinite sequence of numbered commands, also called log entries. The entries are numbered with index numbers. The log entry with a given index will eventually be committed. At that point, your Raft should send the log entry to the larger service for it to execute.

Note: Only RPC may be used for interaction between different Raft instances. For example, different instances of your Raft implementation are not allowed to share Go variables. Your implementation should not use files at all.

In this lab you'll implement most of the Raft design described in the extended paper describing it, including saving persistent state and reading it after a node fails and then restarts. You will not implement cluster membership changes (Section 6) or log compaction / snapshotting (Section 7).

This is a list of useful resources:

- Raft extended paper.
- Lecture notes uploaded to the course piazza page.
- Advice offered to students at MIT in previous offerings of the course.
- Illustrated guide of Raft.

Additional resources: Paxos, Chubby, Paxos Made Live, Spanner, Zookeeper, Harp, Viewstamped Replication, and Paxos Replicated State Machine.

Hints:

- Start early. Although the amount of code to implement isn't large, getting it to work correctly will be very challenging. Both the algorithm and the code is tricky and there are many corner cases to consider. When one of the tests fails, it may take a bit of puzzling to understand in what scenario your solution isn't correct, and how to fix your solution.
- Read and understand the extended Raft paper and the Raft lecture notes before you start. Your implementation should follow the paper's description closely, particularly Figure 2 in the paper, since that's what the tests expect.

Software

Do a `git pull` to get the latest lab software. We supply you with skeleton code and tests in `src/c/raft`, and a simple RPC-like system in `src/labrpc`.

To get up and running, execute the following commands:

```
$ cd golabs-S2017
$ pull
$ export "GOPATH=$PWD"
$ cd "$GOPATH/src/raft"
$ go test
Test (3A): initial election ...
— FAIL: TestInitialElection (5.03s)
config.go:270: expected one leader, got none
Test (3A): election after network failure ...
— FAIL: TestReElection (5.03s)
config.go:270: expected one leader, got none
...
$
```

When you've finished all three parts of the lab, your implementation should pass all the tests in the `src/raft` directory:

```
$ go test
Test (3A): initial election ...
... Passed
Test (3A): election after network failure ...
... Passed
Test (3B): basic agreement ...
... Passed
...
PASS
ok raft 14.587s

$
```

The Code

Implement Raft by adding code to `raft/raft.go`. In that file you'll find a bit of skeleton code, plus examples of how to send and receive RPCs.

Your implementation must support the following interface, which the tester and (eventually) your key/value server will use. You'll find more details in comments in `raft.go`.

```
// create a new Raft server instance:
rf := Make(peers, me, persister, applyCh)

// start agreement on a new log entry:
rf.Start(command interface{}) (index, term, isleader)

// ask a Raft for its current term, and whether it thinks it is leader
rf.GetState() (term, isLeader)

// each time a new entry is committed to the log, each Raft peer
// should send an ApplyMsg to the service (or tester).
type ApplyMsg
```

A service calls `Make(peers, me,)` to create a Raft peer. The `peers` argument is an array of established RPC connections, one to each Raft peer (including this one). The `me` argument is the index of this peer in the `peers` array. `Start(command)` asks Raft to start the processing to append the command to the replicated log. `Start()` should return immediately, without waiting for this process to complete. The service expects your implementation to send an `ApplyMsg` for each new committed log entry to the `applyCh` argument to `Make()`.

Your Raft peers should exchange RPCs using the `labrpc` Go package that we provide to you. It is modeled after Go's `rpc` library, but internally uses Go channels rather than sockets. `raft.go` contains some example code that sends an RPC (`sendRequestVote()`) and that handles an incoming RPC (`RequestVote()`). The reason you must use `labrpc` instead of Go's `RPC` package is that the tester tells `labrpc` to delay RPCs, re-order them, and delete them to simulate challenging network conditions under which your code should work correctly. Don't modify `labrpc` because

we will test your code with the `labrpc` as handed out.

Note: This lab may be your first exposure to writing challenging concurrent code and your first implementation may not be clean enough that you can easily reason about its correctness. Give yourself enough time to rewrite your implementation so that you can easily reason about its correctness. Subsequent labs will build on this lab, so it is important to do a good job on your implementation.

Part 3A

Task: Implement leader election and heartbeats (`AppendEntries` RPCs with no log entries). The goal for Part 3A is for a single leader to be elected, for the leader to remain the leader if there are no failures, and for a new leader to take over if the old leader fails or if packets to/from the old leader are lost. Run `go test -run 3A` to test your 3A code. You also need to make sure that your code is free of races.

Your code should follow Figure 2 in the Raft paper as closely as possible. You will need to add code as follows:

- Add any state you need to the `Raft` struct in `raft.go`. You'll also need to define a struct to hold information about each log entry.
- Modify `Make()` to initialize the states in the `Raft` struct upon the creation of a peer (Raft server).
- Fill in the `RequestVoteArgs` and `RequestVoteReply` structs.
- Implement the `RequestVote()` RPC handler so that servers will vote for one another.
- Modify `Make()` to create a background goroutine that will kick off leader election periodically by sending out `RequestVote` RPCs when it hasn't heard from another peer for a while. This way a peer will learn who is the leader, if there is already a leader, or become the leader itself.
- To implement heartbeats, define an `AppendEntries` RPC struct (though you may not need all the arguments yet), and have the leader send them out periodically. Write an `AppendEntries` RPC handler method that resets the election timeout so that other servers don't step forward as leaders when one has already been elected.

Hints:

- Go RPC sends only struct fields whose names start with capital letters. Sub-structures must also have capitalized field names (e.g. fields of log records in an array). Forgetting to capitalize field names sent by RPC is the single most frequent source of bugs in these labs.
- Make sure the election timeouts in different peers don't always fire at the same time, or else all peers will vote only for themselves and no one will become the leader.
- The tester requires that the leader send heartbeat RPCs no more than ten times per second.

- The tester requires your Raft to elect a new leader within five seconds of the failure of the old leader (if a majority of peers can still communicate). Remember, however, that leader election may require multiple rounds in case of a split vote (which can happen if packets are lost or if candidates unluckily choose the same random backoff times). You must pick election timeouts (and thus heartbeat intervals) that are short enough that it's very likely that an election will complete in less than five seconds even if it requires multiple rounds.
- The paper's Section 5.2 mentions election timeouts in the range of 150 to 300 milliseconds. Such a range only makes sense if the leader sends heartbeats considerably more often than once per 150 milliseconds. Because the tester limits you to 10 heartbeats per second, you will have to use an election timeout larger than the paper's 150 to 300 milliseconds, but not too large, because then you may fail to elect a leader within five seconds.
- You may find Go's `time` and `rand` packages useful. You can also consult `raft/config.go` for an example using `time`.
- If your code has trouble passing the tests, read the paper's Figure 2 again; the full logic for leader election is spread over multiple parts of the figure.
- A good way to debug your code is to insert print statements when a peer sends or receives a message, and collect the output in a file with `go test -run 3A > out`. Then, by studying the trace of messages in the out file, you can identify where your implementation deviates from the desired protocol. You might find `DPrintf` in `util.go` useful to turn printing on and off as you debug different problems.
- You should check your code with `go test -race -run 3A`, and fix any races it reports.

Be sure you pass the 3A tests before submitting Part 3A. Note that the 3A tests test the basic operation of leader election. Parts B and C will test leader election in more challenging settings and may expose bugs in your leader election code which the 3A tests miss.

Important:

Before submitting, please run the 3A tests one final time. Some bugs may not appear on every run, so run the tests multiple times. Make sure that you have addressed all races.

Part 3B

We want Raft to keep a consistent, replicated log of operations. A call to `Start()` at the leader starts the process of adding a new operation to the log; the leader sends the new operation to the other servers in `AppendEntries` RPCs.

Task: Implement the leader and follower code to append new log entries. This will involve implementing `Start()`, completing the `AppendEntries` RPC structs, sending them, fleshing out the `AppendEntry` RPC handler, and advancing the `commitIndex` at the leader. Your first goal should be to pass the `TestBasicAgree()` test (in `test_test.go`). Once you have that working, you should get all the 3B tests to pass (`go test -run 3B`).

You will need to add code as follows:

- You will need to implement the election restriction (section 5.4.1 in the paper).
- Implement the logic needed for `Start()`.
- Update `AppendEntries` RPC structs and RPC handler to support updating log entries.

Hints:

- While the Raft leader is the only server that initiates appends of new entries to the log, all the servers need to independently give each newly committed entry to their local service replica (via their own `applyCh`). You should try to keep the goroutines that implement the Raft protocol as separate as possible from the code that sends committed log entries on the `applyCh` (e.g., by using a separate goroutine for delivering committed messages). If you don't separate these activities cleanly, then it is easy to create deadlocks, either in this lab or in subsequent labs in which you implement services that use your Raft package. Without a clean separation, a common deadlock scenario is as follows: an RPC handler sends on the `applyCh` but it blocks because no goroutine is reading from the channel (e.g., perhaps because it called `Start()`). Now, the RPC handler is blocked while holding the mutex on the Raft structure. The reading goroutine is also blocked on the mutex because `Start()` needs to acquire it. Furthermore, no other RPC handler that needs the lock on the Raft structure can run.
- Give yourself enough time to rewrite your implementation because only after writing a first implementation will you realize how to organize your code cleanly. For example, only after writing one implementation will you understand how to write an implementation that makes it easy to argue that your implementation has no deadlocks.
- Figure out the minimum number of messages Raft should use when reaching agreement in non-failure cases and make your implementation use that minimum.

Important:

Before submitting, please run the 3A and 3B tests one final time. Some bugs may not appear on every run, so run the tests multiple times. Make sure that your code is free of races.

Part 3C

If a Raft-based server reboots it should resume service where it left off. This requires that Raft keep persistent state that survives a reboot. The paper's Figure 2 mentions which state should be persistent, and `raft.go` contains examples of how to save and restore persistent state.

A real implementation would do this by writing Raft's persistent state to disk each time it changes, and reading the latest saved state from disk when restarting after a reboot. Your implementation won't use the disk; instead, it will save and restore persistent state from a `Persister` object (see `persister.go`). Whoever calls `Raft.Make()` supplies a `Persister` that initially holds Raft's most recently persisted state (if any). Raft should initialize its state from that `Persister`, and should use it to save its persistent state each time the state changes. Use the `Persister`'s `ReadRaftState()` and `SaveRaftState()` methods.

Task: Implement persistence by first adding code that saves and restores persistent state to `persist()` and `readPersist()` in `raft.go`. You will need to encode (or "serialize") the state as an array of bytes in order to pass it to the `Persister`. Use Go's `gob` encoder to do this; see the comments in `persist()` and `readPersist()`.

Task: You now need to determine at what points in the Raft protocol your servers are required to persist their state, and insert calls to `persist()` in those places. You must also load persisted state in `Raft.Make()`. Once you've done this, you should pass the remaining tests. You may want to first try to pass the "basic persistence" test (`go test -run 'TestPersist13C'`), and then tackle the remaining ones (`go test -run 3C`).

Note: In order to avoid running out of memory, Raft must periodically discard old log entries, but you **do not** have to worry about this until the next lab.

Hints:

- Many of the 3C tests involve servers failing and the network losing RPC requests or replies.
- The Go `gob` encoder you'll use to encode persistent state only saves fields whose names start with upper case letters. Using small caps for field names is a common source of mysterious bugs, since Go doesn't warn you that they won't be saved.
- In order to pass some of the challenging tests towards the end, such as those marked "unreliable", you will need to implement the optimization to allow a follower to back up the leader's nextIndex by more than one entry at a time. See the description in the extended Raft paper starting at the bottom of page 7 and top of page 8 (marked by a gray line). The lectures will have more details.

Important:

Before submitting, please run all the tests one final time. Some bugs may not appear on every run, so run the tests multiple times.

Collaboration Policy

In addition to the logistics we discussed at the first lecture, note the following:

- You must write all the code you hand in for this lab by yourself, except for code given to you as part of the assignment.
- You are not allowed to look at anyone else's solution, and you are not allowed to look at solutions from previous years.
- You are not allowed to look at other Raft implementations.
- You may discuss the assignments with other students, but you may not look at or copy each others' code. The reason for this rule is that we believe you will learn the most by designing and implementing your lab solution code yourself.
- Please do not publish your code or make it available on any public repository, such as `github.com`.
- Do not share your code with current or future students.

Deliverables

- For each part of the lab submit a zipped file called `lab3PartY_XX.zip`, where Y is the part number (A, B, or C) and XX is your ID. This zip file should only contain the files that you have changed in this lab. For example, for part A, it will contain `raft.go`.
- You need to send your code to `CS432S17@gmail.com` by the deadline at 8 AM on the specified dates. The subject line should be: "Lab Assignment 3Y - SID:XX".
- The assignment will be discussed in the lab.